



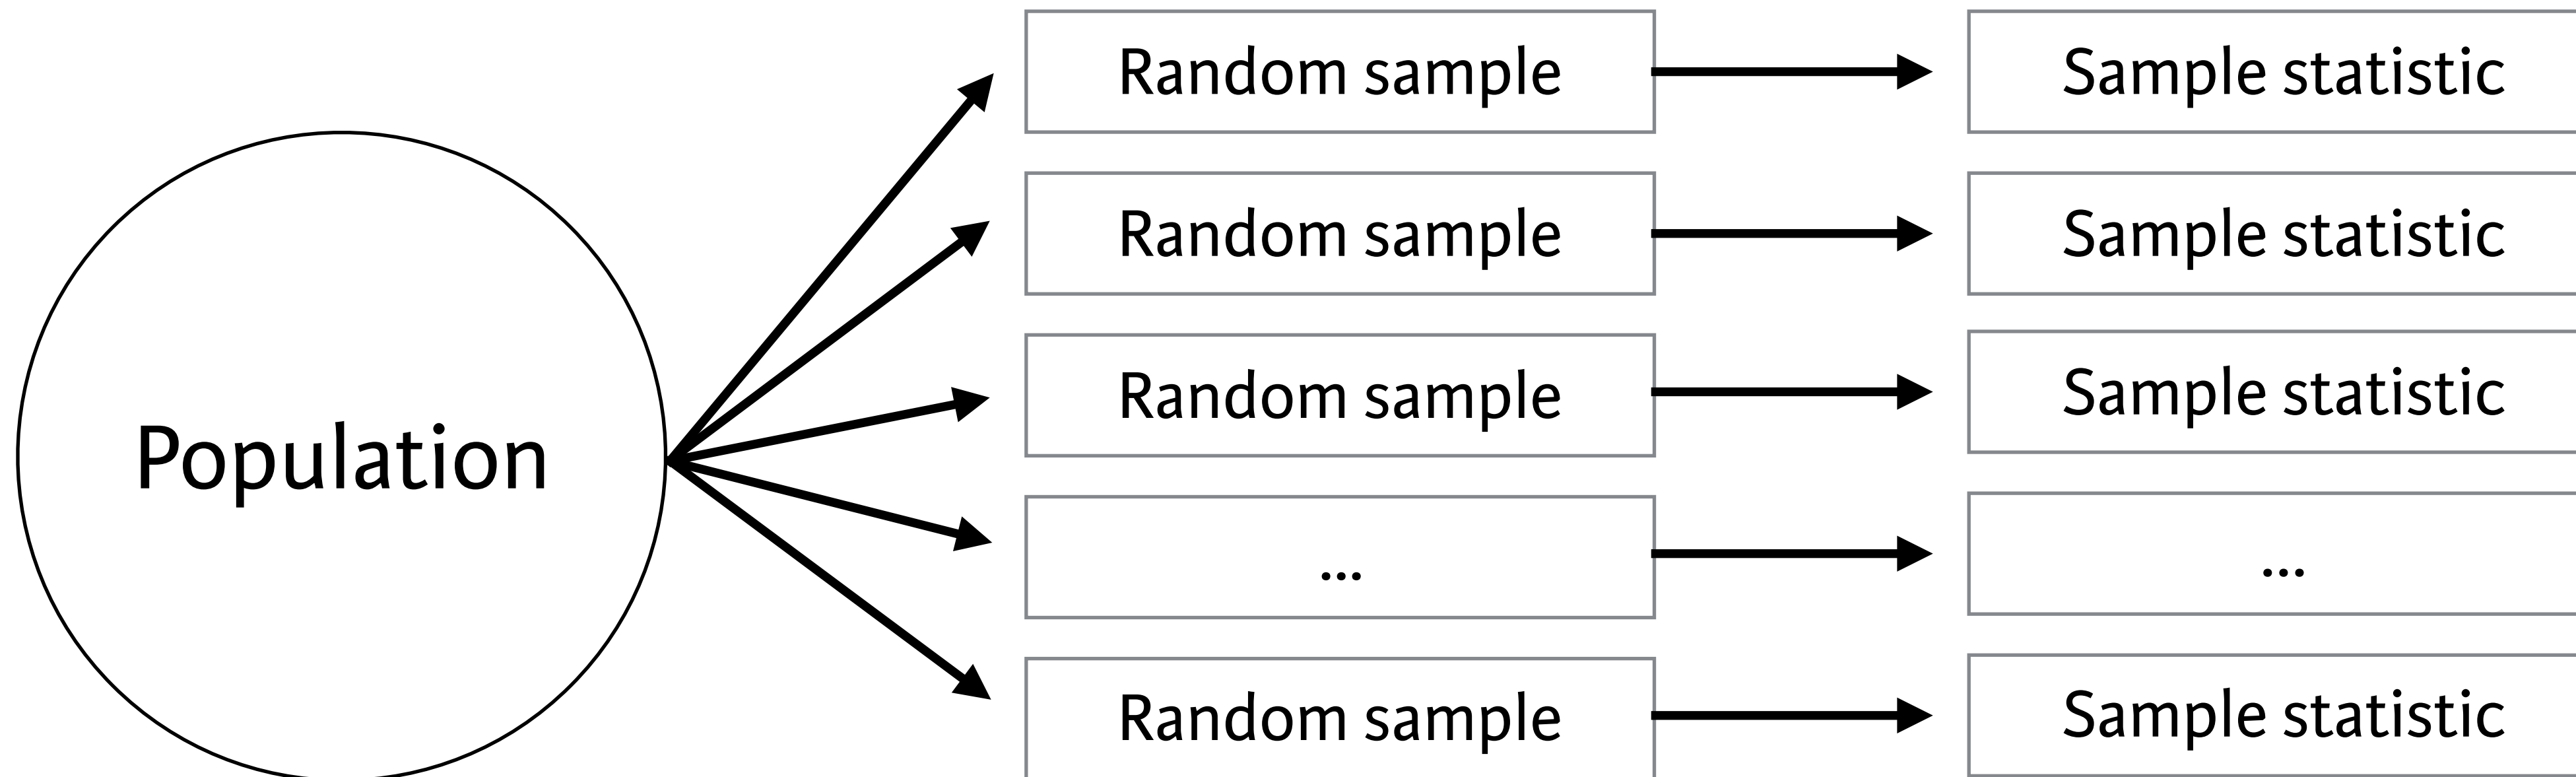
INTRODUCTION TO STATISTICAL MODELING

Precision of effect size estimates

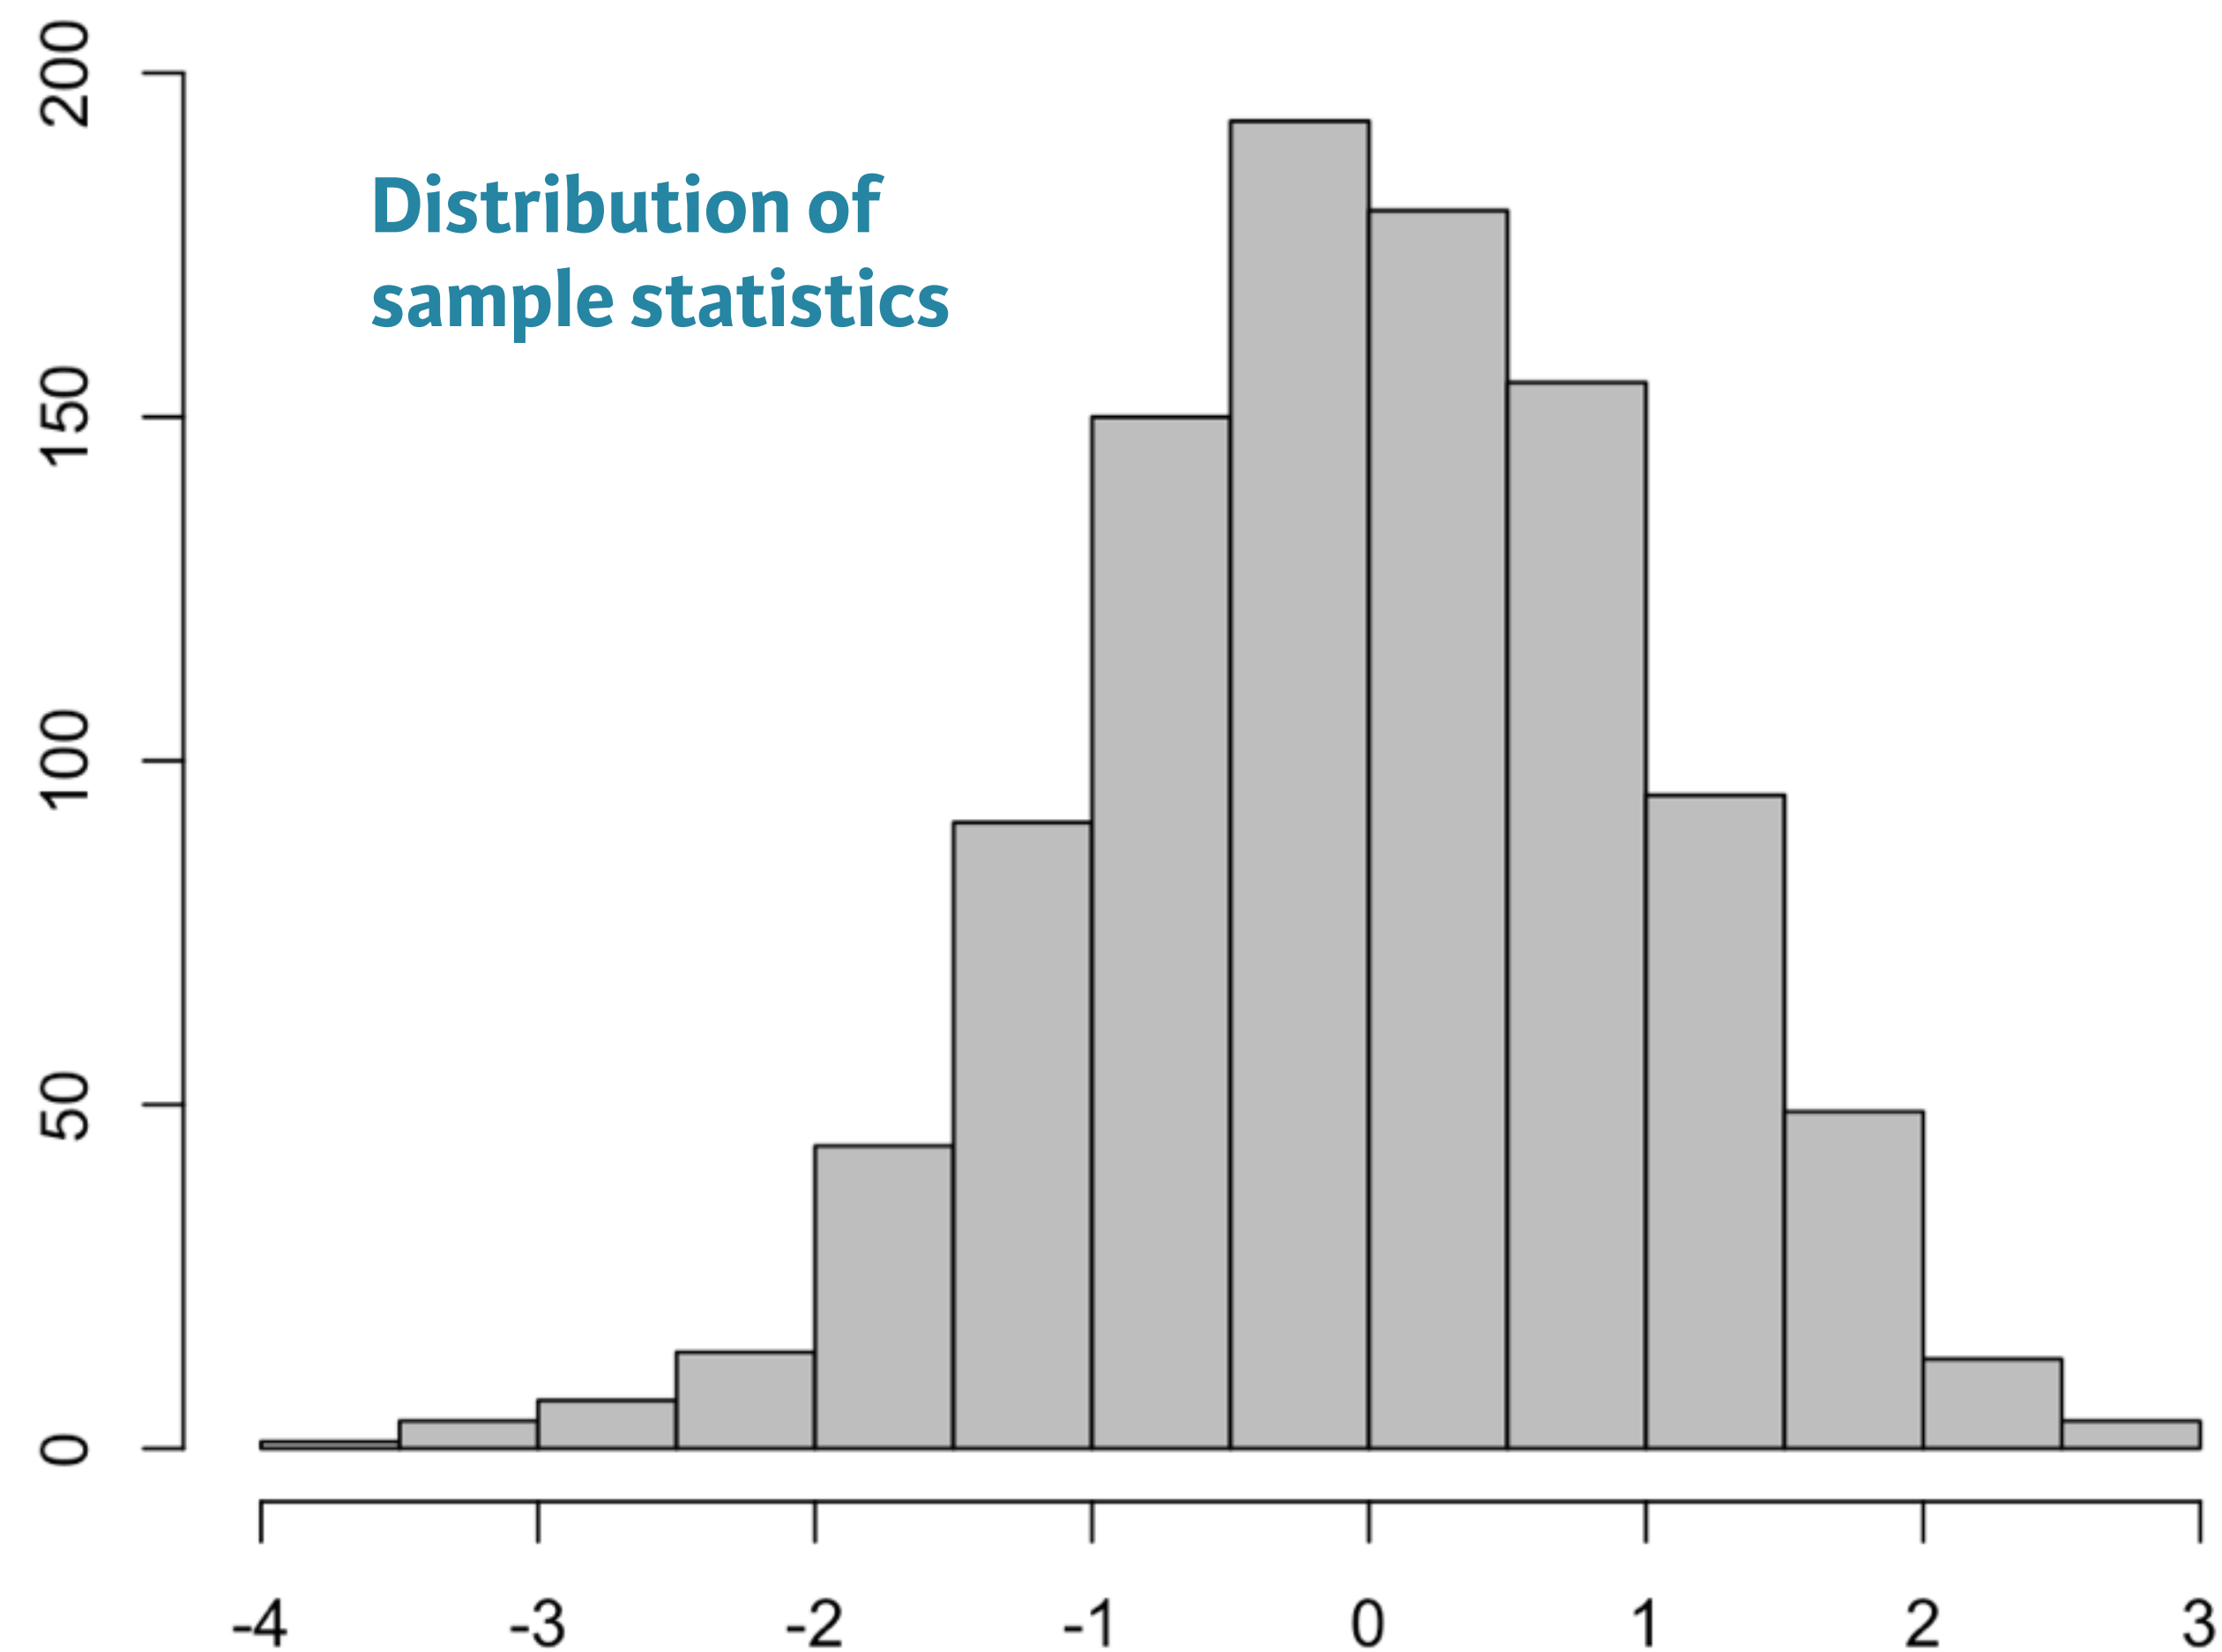
Precision in modeling

- Confidence intervals are a representation of precision
- Take a computational perspective
 - Previous chapters: cross-validation
 - This chapter: bootstrapping
- Precision of effect sizes, prediction error, ...

Framework for thinking about precision

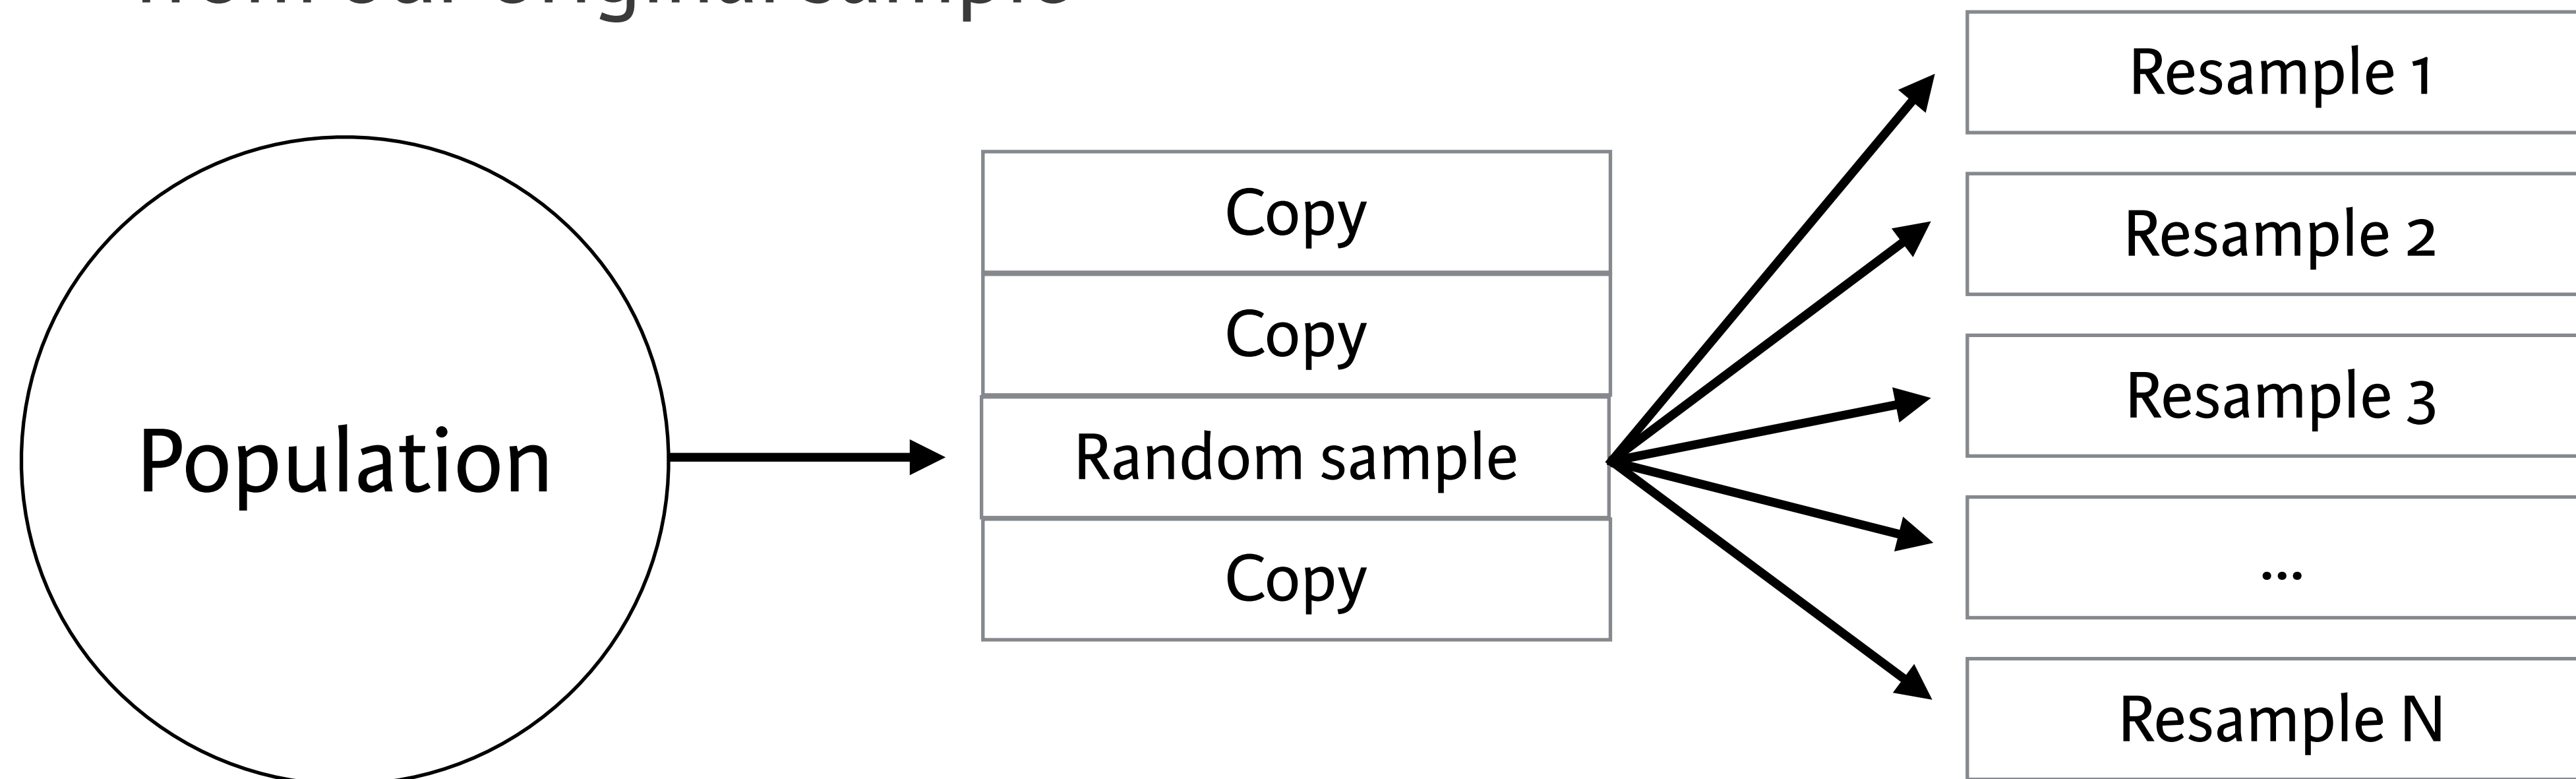


Sampling distribution



Bootstrapping

- Use a single random sample to simulate a situation where you have many random samples
- Use **resampling** to create “new” random samples from our original sample



Resampling in practice

```
> sample_of_statisticians
      name  yob
1      Efron 1938
2 Nightingale 1820
3         Box 1919
```

```
> library(mosaic)

# First resample
> sample(sample_of_statisticians, replace = TRUE)
      name  yob orig.id
1      Efron 1938      1
1.1 Efron 1938      1
3         Box 1919      3

# Second resample
> sample(sample_of_statisticians, replace = TRUE)
      name  yob orig.id
3         Box 1919      3
2 Nightingale 1820      2
3.1         Box 1919      3
```

Rinse and repeat

Resampling and effect size

```
# Train a model on the data
> model <- lm(wage ~ age + exper + sex + sector, data = CPS85)
> effect_size(model, ~ age)
  slope age to:age exper sex sector
1  0.71  35      47    15  M  prof

# Train a model on resampled data
> model_resamp <- lm(wage ~ age + exper + sex + sector,
  data = resample(CPS85))
> effect_size(model_resamp, ~ age)
  slope age to:age exper sex sector
1  0.66  36      48    16  M clerical

# Each new resample is different from the previous ones
> model_resamp <- lm(wage ~ age + exper + sex + sector,
  data = resample(CPS85))
> effect_size(model_resamp, ~ age)
  slope age to:age exper sex sector
1  0.87  34      46    15  M  prof
```



INTRODUCTION TO STATISTICAL MODELING

Let's practice!



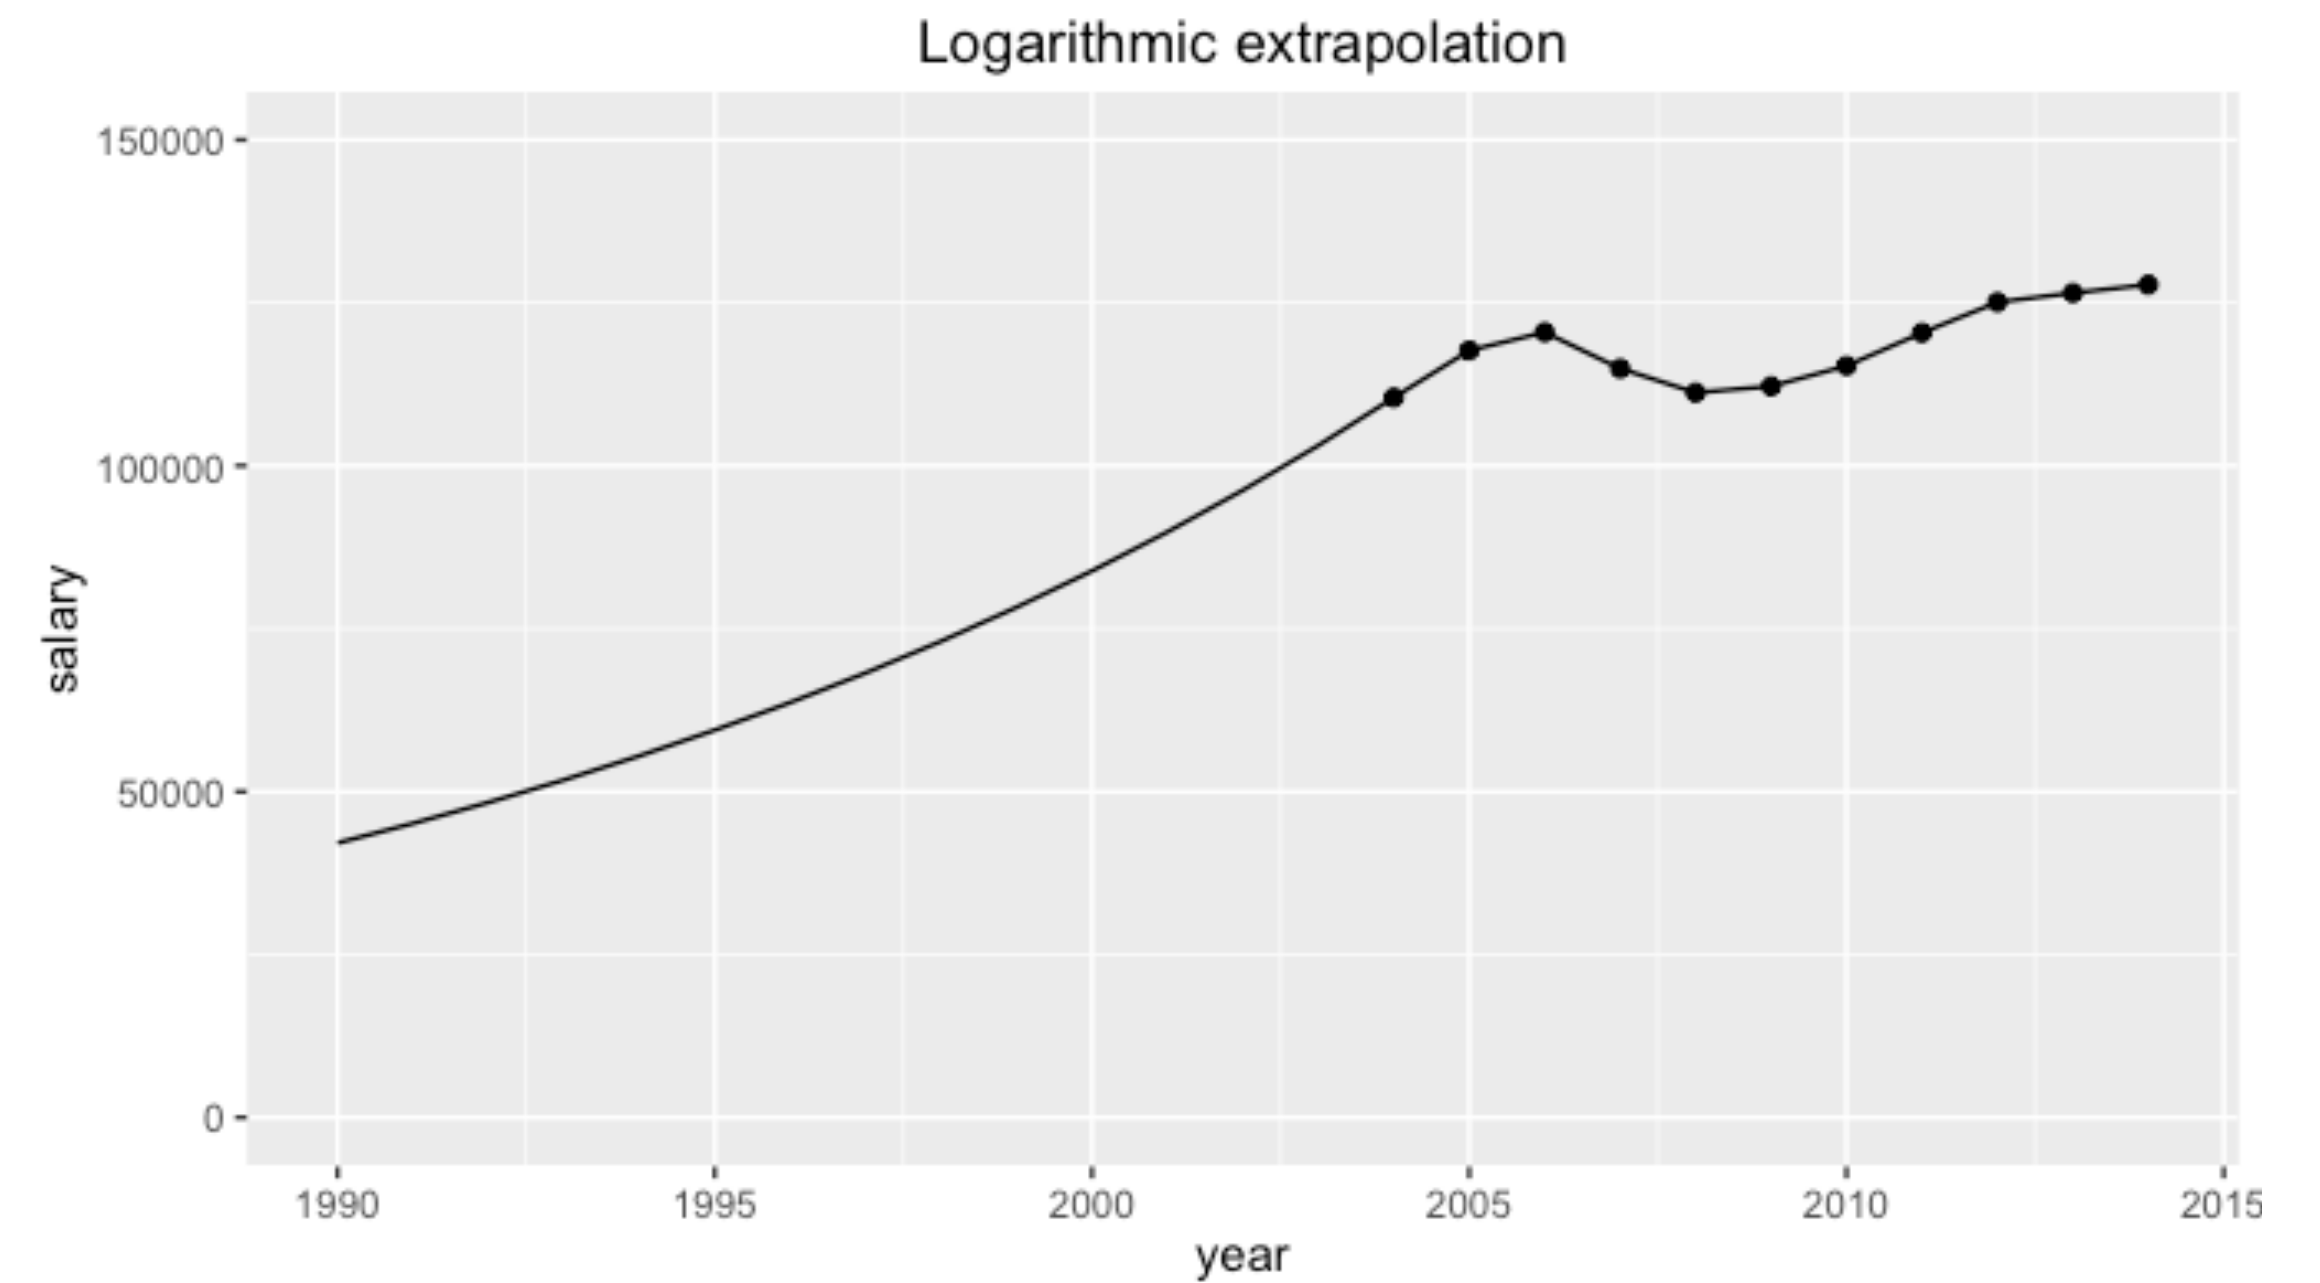

INTRODUCTION TO STATISTICAL MODELING

Scales and transformations

Special situations

- Is the response 0 - 1?
- Is the response a count?
- Is the response cyclic (i.e. day of the week 1-7)?
- Is the response variable money, or another variable where change is proportional to current size?

Modeling salaries



Communicating about logarithms

```
# Effects of age on log wage
> CPS85$log_wage <- log(CPS85$wage)
> model <- lm(log_wage ~ age + exper + sex + sector,
              data = CPS85)

> effect_size(model, ~ age)
      slope age to:age exper sex sector
1 0.075    35      47    15   M   prof

# Convert to a proportion by exponentiation
> exp(0.075) - 1 # subtract 1 for interpretability
[1] 0.078 Wages go up 7.8 percent per year of age
```

The rank transform

The `rank()` function replaces each value with the position that value would occupy if the variable were sorted from least to most


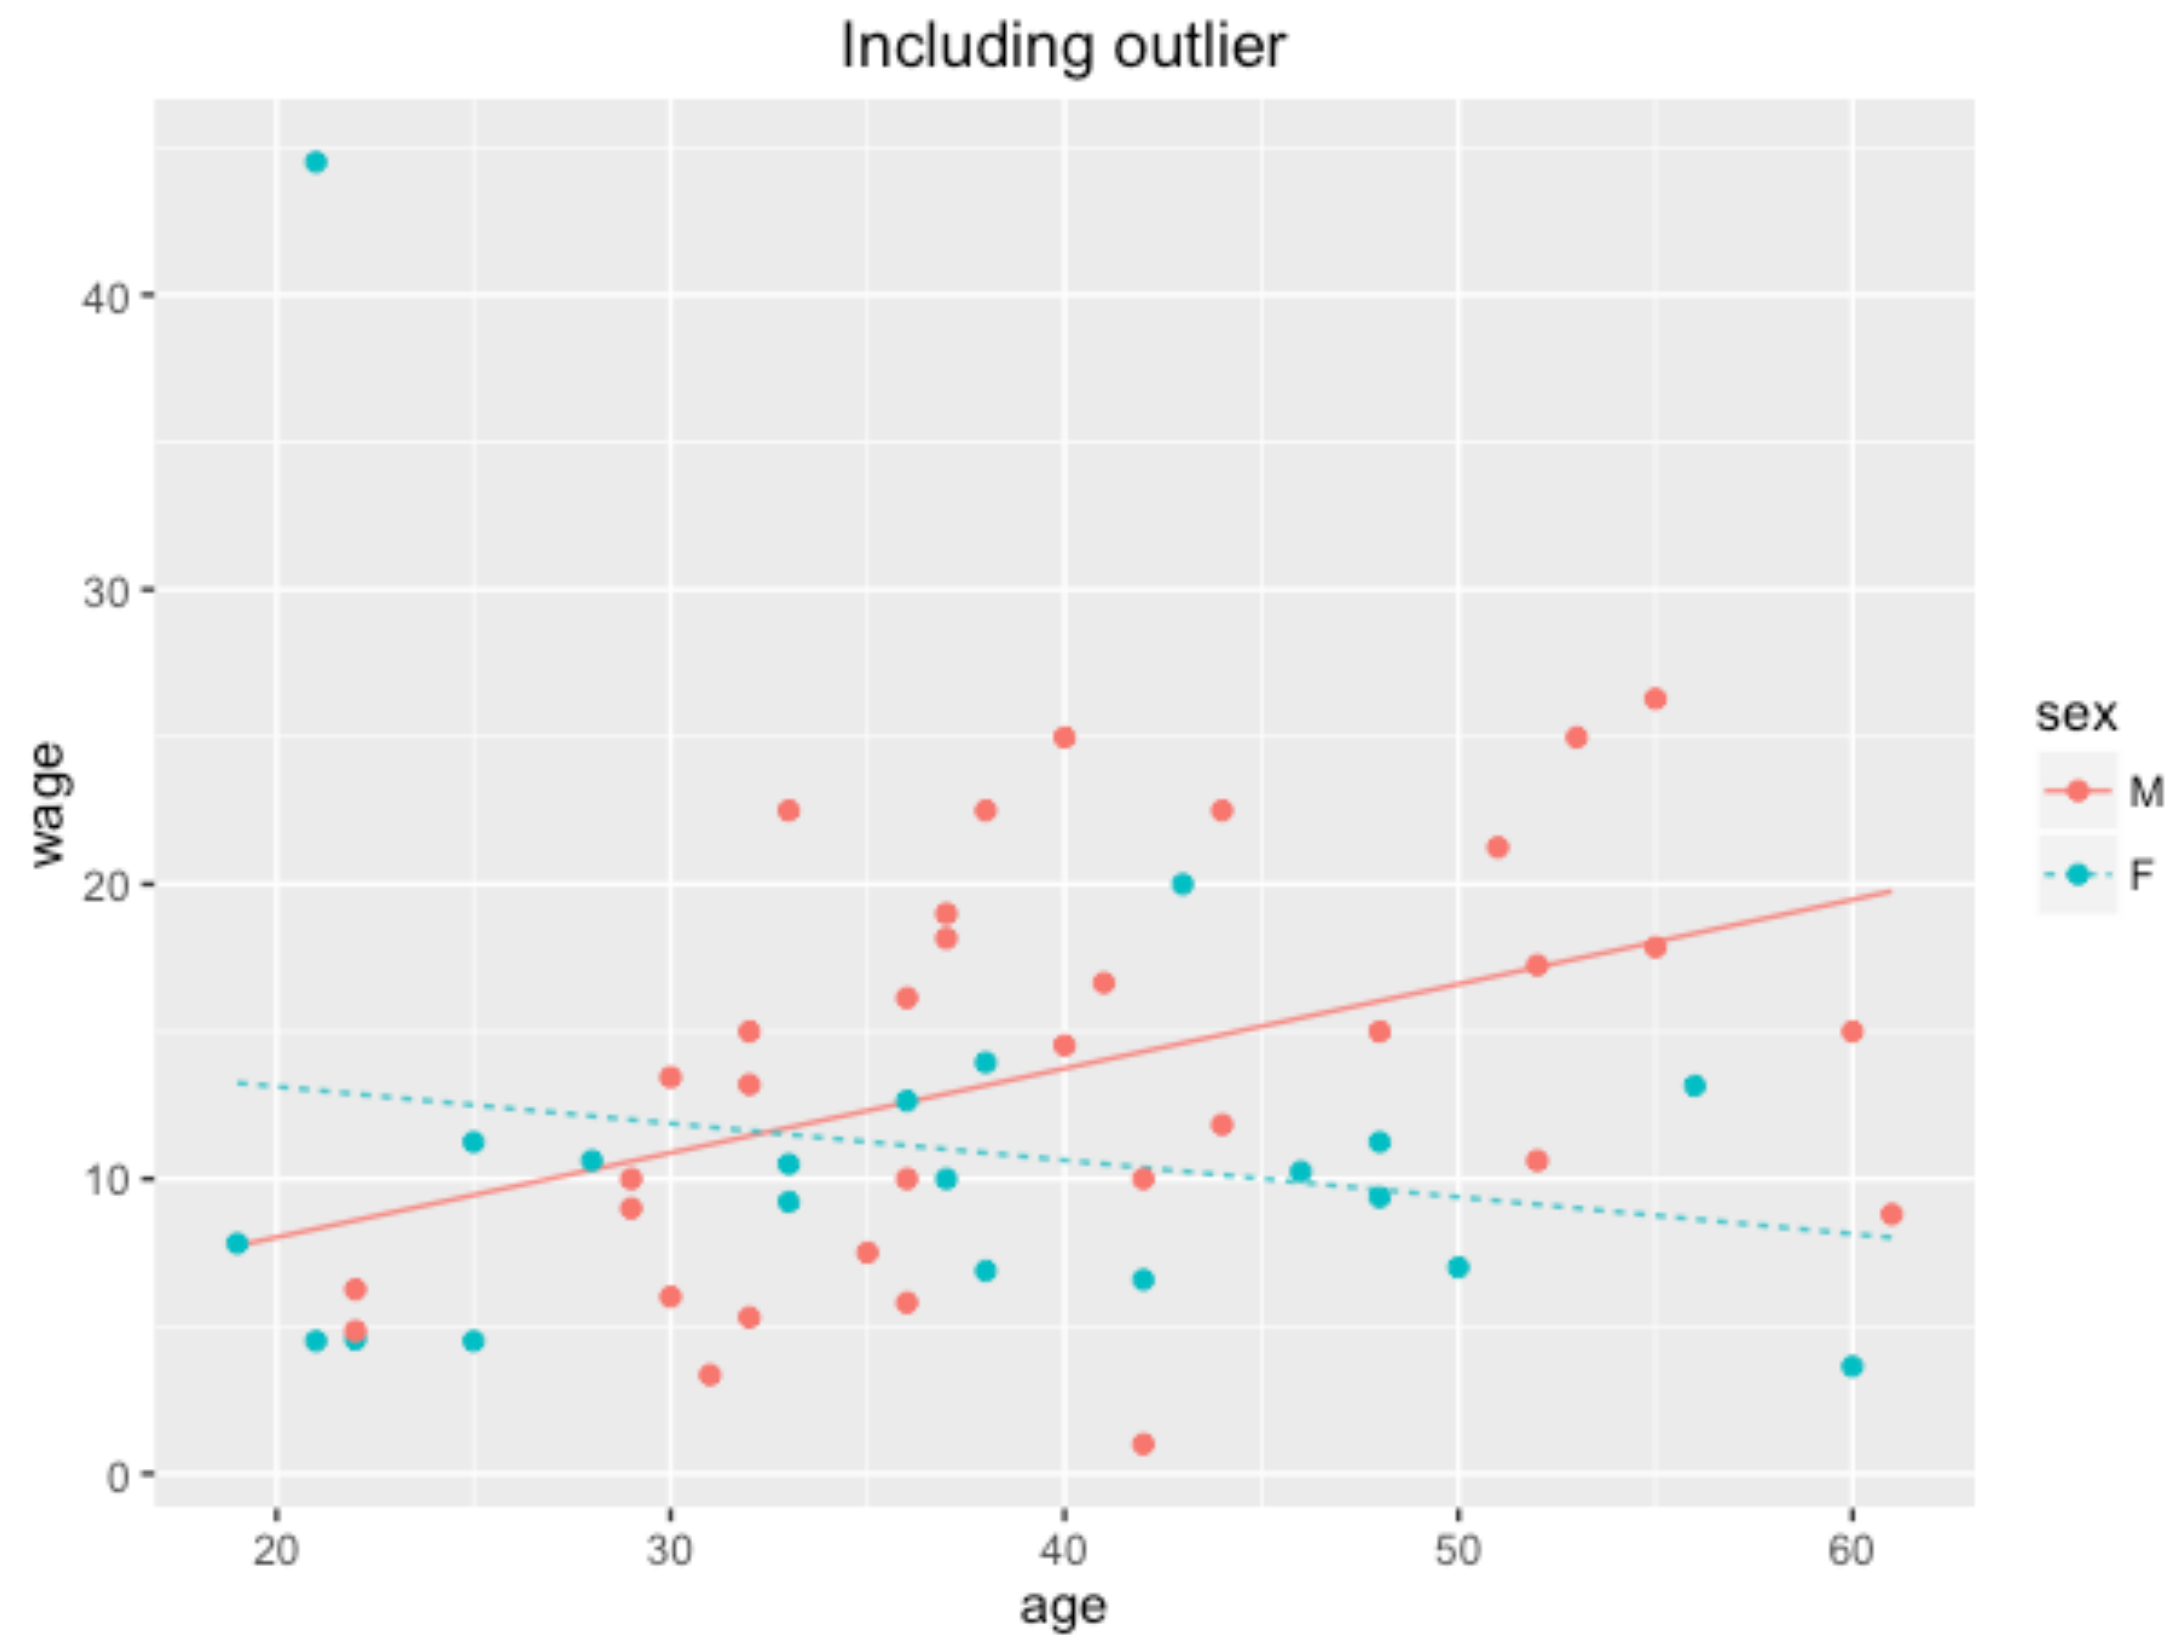
```
> example_vals <- c(2, 1000, 3, 8)
> example_vals
[1] 2 1000 3 8

> rank(example_vals)
[1] 1 4 2 3
```

Spotting an outlier



Applying rank transformation





INTRODUCTION TO STATISTICAL MODELING

Let's practice!